

BAJAJ DCX · RESEARCH PLATFORM

DCX Atlas

One searchable home where DCX's research and customer data outlive the project that produced them.

PART 01 The proposal →

PART 02 The architecture →

Three phases. Research first.

The full scope stands. Macroeconomic and geopolitical tracking is out; everything else is sequenced, not cut.

PHASE 1 • NOW

The research layer

Every study held whole — writeup, links, insights.

Tagged on eight facets, searched by meaning.

Cited answers, evidence one click away.

One owner, one-click gate.

PHASE 2

Internal data

SOR, funnel, clickstream, NPS, complaints, policy. Read on a cadence, not authored.

PHASE 3

Market watch

Competitor launches, pricing, industry shifts — and the chain to their effect inside.

WHY THIS ORDER

The research layer comes first.

WHERE IT BREAKS TODAY

Duplicate work. No one knows the prior research exists.

Scattered & lost. Across folders and inboxes; context leaves with people.

Buried insights. The finding sits inside a deck, never reused.

From interviews across the DT and CX teams.

WHY IT'S THE FOUNDATION

Every phase reads from research.

Only people produce it

A number is read from a database. A finding has to be written.

It's already assumed

The joint architecture names a research store it doesn't have.

Hardest to retrofit

Tagging and evidence must exist before anything can query them.

THE ASSET

Contribute to it. Find from it.

CONTRIBUTE

Submit

By link or file — doc, sheet, Figma.

Highlight → insight

Clip excerpts into tagged atoms.

Guided intake

A short form, product first.

Draft & publish

Save half-done; one-click submit.



THE STUDY ASSET

Writeup

the full study, in text

Links

Figma and recordings

Insights

observation · evidence · tags

AI drafts & suggests — you decide.



FIND

Intent search

Plain language — keyword and meaning.

Cited answer

Every claim links to its source.

Follow-up questions

Ask again on top of any answer.

Related insights

Other findings on the same thing.

Eight facets — six on the study, two on every insight. Lens is already there for phase 2.

Every value we're going with.

PRODUCT · LOANS

Personal Loan
Insta Personal Loan
Loan for Doctors
Loan for CAs
Loan for Lawyers
Business Loan
Home Loan
Loan against Property
Gold Loan
Loan Against Securities
Medical Equipment Finance
Two-wheeler Loan
New Car Finance
Used Car Loan
Loan against Car
Car Loan Balance Transfer and Top-up
New Tractor Loan
Used Tractor Loan
Loan against Tractor
Tractor Loan Balance Transfer

PRODUCT · EMI CARD

Insta EMI Card

PRODUCT · INVESTMENTS

Fixed Deposits
Mutual Funds
Stock Trading

PRODUCT · INSURANCE

Health Insurance
Life Insurance
Car Insurance
Pocket Insurance

PRODUCT · OTHERS

Bajaj Mall
Hamara Mall
Payments

LENS

Customer
Digital
Business
Regulatory
Competition

RESEARCH TYPE

Behavioural
Concept testing
Usability testing
A/B testing
Drop-off
Journey

ALSO ON EVERY STUDY

Business Problem — free text
Date — automatic

OWNER

Sreeraj
Prajakta
Abhishek
Ruchi
Abhijeet

SUBJECT · JOURNEYS

Onboarding
KYC
Application
Transactions
Service
Downloading
Navigation
Calculator
Homepage
CLP
PDP input
PDP product understanding

SUBJECT · CONTENT

Labels
CTA copy
Headings
Body copy
Error messages
Help text
Legal & disclosures
Language / translation

SEGMENT · THE REST

Location — 36 states and UTs
Age — a range
Income — a range

SUBJECT · COMPONENTS

Button / CTA
Card
Input field
Search bar
Checkbox / Radio
Slider
Toggle
Drawer
Tooltip
Banner
Carousel
Accordion
Tabs
Header
Bottom nav
Stepper / Progress
Filter

SEGMENT · LIFECYCLE

NTB — new to Bajaj
PTB — potential to Bajaj
ETB — existing to Bajaj

SEGMENT · REGION

India-1
India-2
India-3

THE BUILD

One app, one database, one external call.

THE STACK

App	Next.js + TypeScript — UI and server in one
Store	PGlite + pgvector — Postgres as a package, no server
Auth	Local login now → BFL SSO in production
AI	Claude Sonnet via OpenRouter · Gemma runs local

ANSWERING A QUESTION

- 1 Embed the question — locally.
- 2 Hybrid search — keyword + meaning — finds the closest insights.
- 3 OpenRouter drafts an answer from only those.
- 4 Real evidence pulled by the cited ID — never invented.

OWN-INFRA · MY MACHINE FOR THE POC, THEN BFL SERVERS

Only the question ever leaves — never the research.

WORKING TODAY

A study goes in. Tagged insights come out.

01

Bring it in

Word, Markdown, or written straight into the editor.

02

Atlas reads it

Title, business problem, product and type, proposed.

03

Insights come out tagged

Every tag from the controlled vocabulary.

04

Embedded on save

Searchable by meaning the moment it's published.

Ninety-two automated checks — the schema, the vocabulary and the vectors.

Trustworthy, and out of the way.

Contributor

Anyone on DT or CX. Writes the whole study — writeup, insights, tags — then clicks submit.

Curator

One named owner. One-click approve, owns the taxonomy, gardens it. Reviewer and owner are one person.

Consumer

Everyone. Searches and reuses — no gate, no friction, ever.

One gate. Submit → one-click accept → live within a day.

One owner. Named from day one; about half a day a week.

Flag, never hide. A 12-month review flag; nothing leaves search.

WHY THE OWNER COMES FIRST

Ownerlessness is the single biggest cause of repository death — so the seat exists before it's even named.

THE ASK

Two months to phase one.

WHAT I'LL DO

1

MONTH 1 • CLOSE THE LOOP

Cited answers

Search, then the studies in.

2

MONTH 2 • MAKE IT THEIRS

On BFL servers

Curator tools, the full set.

WHERE I NEED YOU

A research owner

Named by you, to hold the curator seat.

A tech partner from BFL

Hosts the backend and phase 2's data path.

Sign-off to build

Approval to start — POC, then full build.

PART 02 · THE ARCHITECTURE

How Atlas works.

Part one said what Atlas does. This part opens the machine — how a study gets in, the gate that guards the pool, how search finds and answers, and three real questions run end to end. All of it runs today.

What happens when you click Import.

01 • THE FILE

Word or Markdown

The writeup you already wrote — no special format.

Reads up to 240,000 characters.



02 • THE AI, BRIEFED

Company context

who we are, what we sell

The vocabulary

every tag, with its written definition

The rules

verbatim evidence • closed tags only

Not a blank canvas.



03 • JSON BACK

The study

title • why • when • sample size • study tags

The insights

observation • verbatim evidence • tags or N/A

Checked by code

a fixed schema; every tag must match a real database record



04 • A HUMAN REVIEWS

You decide

The AI proposes; the contributor corrects, deletes, adds.

Nothing is committed by AI alone.

GUARDRAILS ON THE WAY BACK


Verbatim evidence. Lifted from the writeup — never written by the model.

Closed vocabulary. An invented tag is dropped on arrival.

Nominations. A new term becomes a proposal for the curator — never a tag.

~2 minutes a run — 55,000 characters in, 30 tagged insights out.

Only published work can be found.



Checkpoint 0 · the AI’s leash.

Verbatim evidence, closed vocabulary, nothing committed alone.

Checkpoint 1 · the maker.

Corrects every AI proposal, then chooses to submit.

Checkpoint 2 · the checker.

A named curator accepts with one click, or sends it back.


The score. A background pass — code + AI — rates each study’s confidence, with a written why. Only the checker sees it. Mechanism in design.

THE HARD RULE — BUILT IN; THE BUTTON DISABLES

You cannot accept your own work.

Two searches, one honest answer.

Every question runs twice — by meaning and by exact words. No filters, no setup: ask, and the system does the rest.



Two blind searches, one sighted list.

Why two searches

Meaning-search is blind to codes and names it never saw. Word-search is blind to paraphrase. Each sees where the other is blind — so both run, every time.

Combined by position, not score

Each search returns its own ranked list of ~50 insights. Their scores don’t compare — “how close in meaning” and “how strong a word match” are different units — so the combining uses positions. An insight earns points from its rank on each list: high on either earns a place, high on both tops the pile.

Tags describe — they never gate

A study’s tags ride along on every insight, but the ranking never excludes by tag. That is why product X’s research answers product Y’s question.

To the AI

The top ~20 combined insights, with their evidence. It may answer only from these, citing each one it uses.

To the viewer

Every cited insight opens as itself — observation, verbatim evidence, its study and tags, and its confidence marker (scoring in design).

THE SAME QUESTION, THREE INSIGHTS

#2 by meaning • #7 by words

near the top — both searches agree

#4 by meaning • absent by words

still surfaces — a paraphrase words can’t see

absent by meaning • #1 by words

still surfaces — an exact code or name

READING THE ANSWER

Written for whoever's asking.

Same insights, same citations for everyone — the answer only leans toward its reader.

THE DISCIPLINE — FOUR WAYS TO LEAN

Researcher

Evidence quality — method, sample, corroborated or thin.

Designer

The behaviours behind each insight — surfaces, components, flows.

Product

The whole picture — what IT, business and design each need to hear.

Business

Numbers first — conversion, drop-offs — and the behaviour behind them.

One steering sentence in the prompt — never a filter, never a smaller answer.

THE FOLLOW-UP THREAD

Every answer opens a thread. Follow-ups answer from the same retrieved set — the ground never shifts under you.

DEEP RESEARCH — ON DEMAND

Toggled on, Atlas first reads the full writeups of the top 3 studies behind the answer, then answers with that depth. It never adds new sources — the answer stays auditable.

THREE REAL QUESTIONS

Ask it. Watch what comes back.

A · CROSS-PRODUCT

A Gold Loan owner asks — “Do customers trust our interest-rate displays?”

A/B — CONTENT & CTA COPY · TAGGED TWO-WHEELER LOAN

“EMI up to 18%” read as a rate being withheld; “EMI at 18%” lifted starts 26% at 98% confidence.

The tag says Two-wheeler. The meaning matches — RAG surfaces it, because tags were never part of the ranking.

**Tags narrow when you ask them to.
Meaning always searches everything.**

B · BEFORE COMMISSIONING

“Has anyone already studied how customers find a branch?”

LOCATE & FIND · BEHAVIOURAL · JOURNEY

30 tagged insights — 90% open the app to resolve something specific, not to browse.

The study exists. Its insights come back before a brief is written.

Weeks and a budget saved — ask before you commission.

C · ACROSS METHODS

“What frustrates customers about banners?”

STAYING INFORMED · INTERVIEWS

Promotional and account banners look identical — users dismiss both.

A/B — COMPONENT PLACEMENT

Hiding the banner lifted Fixed Deposit starts 15.9%.

Interviews and experiments — one cited answer.

Everything you just saw runs today.

The intake that reads a study, the gate that guards the pool, the search that answers with citations — one system, running end to end. Ask it something.

PART 01 **Back to the ask** →